

Pattern Recognition and Machine Learning (EE602L)

Lab 5: Decision Trees for Classification and Regression

Roll No: EC24BT033

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris, load_diabetes
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor,
plot_tree
```

Question 1: Decision Tree Classification (Iris)

(a) Data preparation and tree training

Iris: 150 flowers, 4 features (sepal/petal length and width), 3 species of 50 each. We reuse the stratified 70:15:15 split from Lab 4 (each class split separately so all sets keep the balanced class mix), seed 42.

Note on scaling: trees need **no standardization**. A tree splits one feature at a time with a threshold ("petal length ≤ 2.45 ?"), and any monotonic rescaling of a feature just rescales its thresholds without changing which points go left or right, so the learned tree is identical. This is one practical advantage of trees over the distance- and gradient-based models of Labs 2-4.

```
In [2]: def stratified_split(df, label_col, train_ratio=0.7, val_ratio=0.15,
test_ratio=0.15, seed=42):
    """Split each class separately so all three sets keep the class
    proportions."""
    assert abs(train_ratio + val_ratio + test_ratio - 1.0) < 1e-9, "ratios
    must sum to 1"
    rng = np.random.default_rng(seed)

    train_parts = []
    val_parts = []
    test_parts = []
    for cls in sorted(df[label_col].unique()):
        cls_idx = df.index[df[label_col] == cls].to_numpy()
        cls_idx = cls_idx[rng.permutation(len(cls_idx))]
        n_train = int(train_ratio * len(cls_idx))
        n_val = int(val_ratio * len(cls_idx))
```

```

train_parts.append(cls_idx[:n_train])
val_parts.append(cls_idx[n_train:n_train + n_val])
test_parts.append(cls_idx[n_train + n_val:])

train_idx = np.concatenate(train_parts)
val_idx = np.concatenate(val_parts)
test_idx = np.concatenate(test_parts)
return df.loc[train_idx], df.loc[val_idx], df.loc[test_idx]

iris = load_iris()
idf = pd.DataFrame(iris.data, columns=[str(c) for c in iris.feature_names])
idf["target"] = iris.target

itrain, ival, itest = stratified_split(idf, "target")
print(f"train {len(itrain)}, val {len(ival)}, test {len(itest)}")

tr, va, te = set(itrain.index), set(ival.index), set(itest.index)
assert not (tr & va) and not (tr & te) and not (va & te), "LEAKAGE: splits overlap!"
assert len(tr | va | te) == len(idf), "rows lost or duplicated"
print("verified: splits disjoint, cover all", len(idf), "rows")
print("class counts per split:")
for name, part in [("train", itrain), ("val", ival), ("test", itest)]:
    print(f" {name}:")
    {part['target'].value_counts().sort_index().tolist()}")

```

```

train 105, val 21, test 24
verified: splits disjoint, cover all 150 rows
class counts per split:
  train: [35, 35, 35]
  val: [7, 7, 7]
  test: [8, 8, 8]

```

```

In [3]: iris_features = [str(c) for c in iris.feature_names]
Xi_train, yi_train = itrain[iris_features], itrain["target"]
Xi_val, yi_val = ival[iris_features], ival["target"]
Xi_test, yi_test = itest[iris_features], itest["target"]

depths = [1, 2, 3, 4, 5, 8, None]      # None = unrestricted

rows = []
for criterion in ["entropy", "gini"]:
    for depth in depths:
        clf = DecisionTreeClassifier(criterion=criterion, max_depth=depth,
random_state=42)
        clf.fit(Xi_train, yi_train)
        rows.append({
            "criterion": criterion,
            "max_depth": "None" if depth is None else depth,
            "actual depth": clf.get_depth(),
            "train acc": clf.score(Xi_train, yi_train),
            "val acc": clf.score(Xi_val, yi_val),
        })

```


```
sweep_q1 = pd.DataFrame(rows)
print(sweep_q1.round(4).to_string(index=False))
```

criterion	max_depth	actual depth	train acc	val acc
entropy	1	1	0.6667	0.6667
entropy	2	2	0.9524	0.9048
entropy	3	3	0.9619	0.9048
entropy	4	4	0.9810	0.8095
entropy	5	5	0.9905	0.8571
entropy	8	6	1.0000	0.8095
entropy	None	6	1.0000	0.8095
gini	1	1	0.6667	0.6667
gini	2	2	0.9524	0.9048
gini	3	3	0.9619	0.9524
gini	4	4	0.9810	0.9524
gini	5	5	1.0000	0.9524
gini	8	5	1.0000	0.9524
gini	None	5	1.0000	0.9524

(b) Overfitting and pruning

```
In [4]: fig, axes = plt.subplots(1, 2, figsize=(14, 5), sharey=True)
depth_axis = [d if d is not None else 12 for d in depths] # plot None as depth 12

for ax, criterion in zip(axes, ["entropy", "gini"]):
    sub = sweep_q1[sweep_q1["criterion"] == criterion]
    ax.plot(depth_axis, sub["train acc"], marker="o", label="train accuracy")
    ax.plot(depth_axis, sub["val acc"], marker="o", label="validation accuracy")
    ax.set_title(f"criterion = {criterion}")
    ax.set_xlabel("max_depth (12 = unrestricted)")
    ax.set_ylabel("accuracy")
    ax.legend()
plt.tight_layout()
plt.show()
```



```
In [5]: # Pick the better criterion and best depth by validation accuracy (ties -> shallower tree)
best = None
for r in rows:
```

```

    if best is None or r["val acc"] > best["val acc"]:
        best = r
    best_criterion = best["criterion"]
    best_depth = None if best["max_depth"] == "None" else
    int(best["max_depth"])
    print(f"best by validation accuracy: criterion = {best_criterion},
    max_depth = {best['max_depth']} (val acc {best['val acc']:.4f})")

    pre_pruned = DecisionTreeClassifier(criterion=best_criterion,
    max_depth=best_depth, random_state=42)
    pre_pruned.fit(Xi_train, yi_train)

```

best by validation accuracy: criterion = gini, max_depth = 3 (val acc 0.9524)

Out[5]:

```

▼ DecisionTreeClassifier ⓘ ?
  ► Parameters
  ► Fitted attributes

```

Where overfitting begins: training accuracy keeps climbing to 100% as depth grows, while validation accuracy plateaus around depth 2-3 and then stays flat or dips. Beyond depth ~3 the extra splits only memorize individual training flowers (leaves with 1-2 samples), which is where overfitting starts. We therefore select the shallow depth that already maximizes validation accuracy.

Post-pruning with cost-complexity (ccp_alpha): grow the unrestricted tree first, then prune back the branches whose impurity improvement does not justify their complexity. ccp_alpha is the price per leaf: alpha = 0 keeps the full tree, larger alpha prunes harder. We take the candidate alphas from sklearn's cost_complexity_pruning_path (the exact thresholds where subtrees collapse) and choose by validation accuracy.

```

In [6]: full_tree = DecisionTreeClassifier(criterion=best_criterion,
    random_state=42)
    full_tree.fit(Xi_train, yi_train)

    path = full_tree.cost_complexity_pruning_path(Xi_train, yi_train)
    alphas = [float(round(a, 5)) for a in path.ccp_alphas]
    print("pruning-path alphas:", alphas)

    post_rows = []
    for alpha in path.ccp_alphas:
        clf = DecisionTreeClassifier(criterion=best_criterion, ccp_alpha=alpha,
        random_state=42)
        clf.fit(Xi_train, yi_train)
        post_rows.append({
            "ccp_alpha": alpha,                                # keep the exact value; round
only for display
            "leaves": clf.get_n_leaves(),
            "depth": clf.get_depth(),
            "train acc": clf.score(Xi_train, yi_train),

```

```

        "val acc": clf.score(Xi_val, yi_val),
    })

post_table = pd.DataFrame(post_rows)
print(post_table.round(4).to_string(index=False))

# Several alphas tie on validation accuracy; ties go to the LARGEST alpha
# (the most pruned tree), consistent with our Lab 4 tie-break: when
validation
# cannot distinguish, keep the simplest model.
best_val = max(r["val acc"] for r in post_rows)
tied = [r for r in post_rows if r["val acc"] == best_val]
best_alpha = max(t["ccp_alpha"] for t in tied)

post_pruned = DecisionTreeClassifier(criterion=best_criterion,
ccp_alpha=best_alpha, random_state=42)
post_pruned.fit(Xi_train, yi_train)
print(f"\nbest ccp_alpha by validation accuracy (ties -> most pruned):
{best_alpha:.5f} "
      f"({post_pruned.get_n_leaves()} leaves, depth
{post_pruned.get_depth()})")

```

pruning-path alphas: [0.0, 0.00923, 0.01429, 0.01845, 0.02115, 0.24671, 0.33333]

ccp_alpha	leaves	depth	train acc	val acc
0.0000	9	5	1.0000	0.9524
0.0092	7	5	0.9905	0.9524
0.0143	5	3	0.9619	0.9524
0.0185	4	3	0.9524	0.9524
0.0211	3	2	0.9524	0.9048
0.2467	2	1	0.6667	0.6667
0.3333	1	0	0.3333	0.3333

best ccp_alpha by validation accuracy (ties -> most pruned): 0.01845 (4 leaves, depth 3)

(c) Evaluation and analysis

```

In [7]: def confusion_matrix_manual(y_true, y_pred, K):
        cm = np.zeros((K, K), dtype=int)
        for t, p in zip(y_true, y_pred):
            cm[t, p] += 1
        return cm

def macro_metrics(cm):
    """Per-class precision/recall/F1 from the confusion matrix, macro-averaged."""
    K = cm.shape[0]
    precisions, recalls, f1s = [], [], []
    for k in range(K):
        tp = cm[k, k]
        fp = cm[:, k].sum() - tp
        fn = cm[k, :].sum() - tp
        prec = tp / (tp + fp) if (tp + fp) > 0 else 0.0
        rec = tp / (tp + fn) if (tp + fn) > 0 else 0.0

```

```

        f1 = 2 * prec * rec / (prec + rec) if (prec + rec) > 0 else 0.0
        precisions.append(prec); recalls.append(rec); f1s.append(f1)
    return np.mean(precisions), np.mean(recalls), np.mean(f1s)

results = []
for name, model in [("pre-pruned (max_depth)", pre_pruned), ("post-pruned
(ccp_alpha)", post_pruned)]:
    y_pred = model.predict(Xi_test)
    cm = confusion_matrix_manual(yi_test.to_numpy(), y_pred, 3)
    prec, rec, f1 = macro_metrics(cm)
    results.append({"model": name, "test accuracy": np.mean(y_pred ==
yi_test.to_numpy()),
                    "macro precision": prec, "macro recall": rec, "macro
F1": f1,
                    "leaves": model.get_n_leaves(), "depth":
model.get_depth()})

print(pd.DataFrame(results).round(4).to_string(index=False))

```


	model	test accuracy	macro precision	macro recall	macro
F1	leaves	depth			
	pre-pruned (max_depth)	1.0	1.0	1.0	
1.0	5	3			
	post-pruned (ccp_alpha)	1.0	1.0	1.0	
1.0	4	3			

```

In [8]: plt.figure(figsize=(14, 8))
        plot_tree(pre_pruned, feature_names=iris_features,
        class_names=list(iris.target_names),
        filled=True, rounded=True, fontsize=9)
        plt.title(f"Selected pre-pruned tree (criterion = {best_criterion},
        max_depth = {pre_pruned.get_depth()})")
        plt.show()

```

Selected pre-pruned tree (criterion = gini, max_depth = 3)



Q1 discussion

Entropy vs Gini: both measure how mixed a node is, and a split is chosen to reduce that mixing the most. Entropy = $-\sum p_k \log_2(p_k)$ comes from information theory (bits of uncertainty); Gini = $1 - \sum p_k^2$ is the chance of mislabeling a random sample if labels were assigned by the node's class proportions. They are numerically very close (Gini is roughly a first-order approximation of entropy), usually pick the same splits, and did essentially equally well here; Gini is marginally cheaper to compute (no logarithm), which is why it is sklearn's default.

Why depth causes overfitting: every extra level doubles the number of possible regions the tree can carve. An unrestricted tree keeps splitting until leaves are pure, which on a small dataset means leaves holding one or two specific training flowers: the tree has memorized samples, not learned species. Training accuracy hits 100% while validation stalls or drops: the gap between those two curves is the overfitting, visible in our plots beyond depth ~3.

Pre-pruning vs post-pruning: pre-pruning (max_depth, min_samples_leaf) stops growth early: cheap, but it must guess the right size before seeing how splits pan out, and a greedy-looking weak split can hide a strong one below it. Post-pruning (ccp_alpha) grows the full tree first and then removes branches whose impurity gain does not pay for their added leaves: it judges each branch by its realized value, at the cost of building the big tree first. Both reached the same test performance here, with the post-pruned tree keeping only the branches worth their price.

Question 2: Decision Tree Regression (Diabetes)

(a) Data preparation and tree training

Diabetes: 442 patients, 10 standardized clinical features (age, sex, BMI, blood pressure, six blood serum measurements), target = a continuous disease-progression score one year after baseline. Continuous target means no stratification (there are no classes to balance); we use the plain shuffled 70:15:15 split from Labs 1-2, seed 42.

```
In [9]: def custom_split(df, train_ratio=0.7, val_ratio=0.15, test_ratio=0.15,
      seed=42):
    """Shuffle and split a dataframe into train/val/test using only
    NumPy/Pandas."""
    assert abs(train_ratio + val_ratio + test_ratio - 1.0) < 1e-9, "ratios
    must sum to 1"
    rng = np.random.default_rng(seed)
    perm = rng.permutation(len(df))
    n_train = int(train_ratio * len(df))
    n_val = int(val_ratio * len(df))
    train_idx = perm[:n_train]
```

```

val_idx = perm[n_train:n_train + n_val]
test_idx = perm[n_train + n_val:]
return df.iloc[train_idx], df.iloc[val_idx], df.iloc[test_idx]

diab = load_diabetes()
ddf = pd.DataFrame(diab.data, columns=[str(c) for c in diab.feature_names])
ddf["target"] = diab.target

dtrain, dval, dtest = custom_split(ddf)
print(f"train {len(dtrain)}, val {len(dval)}, test {len(dtest)}")
tr, va, te = set(dtrain.index), set(dval.index), set(dtest.index)
assert not (tr & va) and not (tr & te) and not (va & te) and len(tr | va | te) == len(ddf)
print("verified: splits disjoint, cover all", len(ddf), "rows")

```

train 309, val 66, test 67

verified: splits disjoint, cover all 442 rows

```

In [10]: diab_features = [str(c) for c in diab.feature_names]
Xd_train, yd_train = dtrain[diab_features], dtrain["target"]
Xd_val, yd_val = dval[diab_features], dval["target"]
Xd_test, yd_test = dtest[diab_features], dtest["target"]

def mse(y, y_hat):
    return np.mean((np.asarray(y) - np.asarray(y_hat)) ** 2)

reg_depths = [1, 2, 3, 4, 5, 8, 10, None]

depth_rows = []
for depth in reg_depths:
    reg = DecisionTreeRegressor(max_depth=depth, random_state=42)
    reg.fit(Xd_train, yd_train)
    depth_rows.append({
        "max_depth": "None" if depth is None else depth,
        "actual depth": reg.get_depth(),
        "train MSE": mse(yd_train, reg.predict(Xd_train)),
        "val MSE": mse(yd_val, reg.predict(Xd_val)),
    })

depth_table = pd.DataFrame(depth_rows)
print(depth_table.round(1).to_string(index=False))

```

max_depth	actual depth	train MSE	val MSE
1	1	4152.3	4245.4
2	2	3467.8	3029.1
3	3	2980.6	2772.0
4	4	2334.4	3181.4
5	5	1704.5	3737.4
8	8	514.5	4351.2
10	10	126.8	5003.3
None	16	0.0	5628.6

(b) Hyperparameter study: min_samples_leaf at fixed max_depth = 10

```
In [11]: leaf_values = [1, 5, 10, 20, 50]

leaf_rows = []
for leaf in leaf_values:
    reg = DecisionTreeRegressor(max_depth=10, min_samples_leaf=leaf,
                                random_state=42)
    reg.fit(Xd_train, yd_train)
    leaf_rows.append({
        "min_samples_leaf": leaf,
        "leaves": reg.get_n_leaves(),
        "train MSE": mse(yd_train, reg.predict(Xd_train)),
        "val MSE": mse(yd_val, reg.predict(Xd_val)),
    })

leaf_table = pd.DataFrame(leaf_rows)
print(leaf_table.round(1).to_string(index=False))
```

min_samples_leaf	leaves	train MSE	val MSE
1	219	126.8	5003.3
5	50	1487.0	3566.1
10	24	2257.8	3551.0
20	11	2944.1	2787.7
50	4	3535.4	3449.6


(c) Evaluation and analysis

```
In [12]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

depth_axis = [d if d is not None else 14 for d in reg_depths]
axes[0].plot(depth_axis, depth_table["train MSE"], marker="o", label="train MSE")
axes[0].plot(depth_axis, depth_table["val MSE"], marker="o",
              label="validation MSE")
axes[0].set_xlabel("max_depth (14 = unrestricted)")
axes[0].set_ylabel("MSE")
axes[0].set_title("MSE vs max_depth (min_samples_leaf = 1)")
axes[0].legend()

axes[1].plot(leaf_values, leaf_table["train MSE"], marker="o", label="train MSE")
axes[1].plot(leaf_values, leaf_table["val MSE"], marker="o",
              label="validation MSE")
axes[1].set_xlabel("min_samples_leaf (max_depth = 10)")
axes[1].set_title("MSE vs min_samples_leaf")
axes[1].legend()

plt.tight_layout()
plt.show()
```



```
In [13]: # Best model by validation MSE across everything we swept
candidates = []
for r in depth_rows:
    d = None if r["max_depth"] == "None" else int(r["max_depth"])
    candidates.append({"max_depth": d, "min_samples_leaf": 1, "val MSE":
r["val MSE"]})
for r in leaf_rows:
    candidates.append({"max_depth": 10, "min_samples_leaf":
r["min_samples_leaf"], "val MSE": r["val MSE"]})

best_cfg = min(candidates, key=lambda c: c["val MSE"])
print("best configuration by validation MSE:", best_cfg)

best_reg = DecisionTreeRegressor(max_depth=best_cfg["max_depth"],
min_samples_leaf=best_cfg["min_samples_leaf"], random_state=42)
best_reg.fit(Xd_train, yd_train)


test_pred = best_reg.predict(Xd_test)
test_mse = mse(yd_test, test_pred)
print(f"\ntest MSE: {test_mse:.1f}")
print(f"test RMSE: {np.sqrt(test_mse):.2f}")
print(f"test MAE: {np.mean(np.abs(yd_test.to_numpy() - test_pred)):.2f}")
print(f"(target ranges {int(yd_test.min())} to {int(yd_test.max())}, leaves
in model: {best_reg.get_n_leaves()})")

best configuration by validation MSE: {'max_depth': 3, 'min_samples_leaf':
1, 'val MSE': np.float64(2771.9805187688603)}

test MSE: 3707.7
test RMSE: 60.89
test MAE: 49.11
(target ranges 31 to 332, leaves in model: 8)
```

```
In [14]: # Visualize the selected tree, truncated to 3 levels for readability
plt.figure(figsize=(16, 8))
plot_tree(best_reg, feature_names=diab_features, max_depth=3, filled=True,
rounded=True, fontsize=8)
plt.title("Selected regression tree (display truncated to depth 3)")
plt.show()
```

Selected regression tree (display truncated to depth 3)



Q2 discussion

How a regression tree picks splits: at each node it tries every feature and threshold, and for each candidate computes the MSE of the two children (each child's MSE measured around its own mean target). It picks the split that reduces total MSE the most versus the parent. It is the same greedy recipe as classification, with variance replacing impurity.

Leaf predictions: a leaf predicts the **mean target value of the training samples that landed in it**. The mean is used because it is exactly the constant that minimizes MSE for that group (Lab 2's lesson: MSE-optimal constant = mean). So the tree's prediction function is a staircase: piecewise-constant regions, one level per leaf, which is also why single regression trees look blocky compared to linear models.

max_depth and min_samples_leaf: the depth sweep shows the classic U: unrestricted depth drives train MSE to 0 (every patient in a private leaf: pure memorization) while validation MSE roughly doubles from its best value; depth 3-4 is the sweet spot. min_samples_leaf attacks the same problem from below: requiring, say, 20 patients per leaf means every prediction is an average over at least 20 people, which caps how specific (and how noisy) a leaf can get. The leaf sweep shows validation MSE improving markedly over leaf = 1 even with depth fixed at 10, because leaf size directly controls the variance of leaf means. The two knobs are complementary pre-pruning controls: depth limits how many questions may be asked, leaf size limits how small an audience an answer may serve.